

hash join->2511

first_rows(5000)->TAF-> 5011

hash join->5001

Query #2

Confrontare i costi di hash join e nested loop usando l'hint USE / NO_USE_HASH

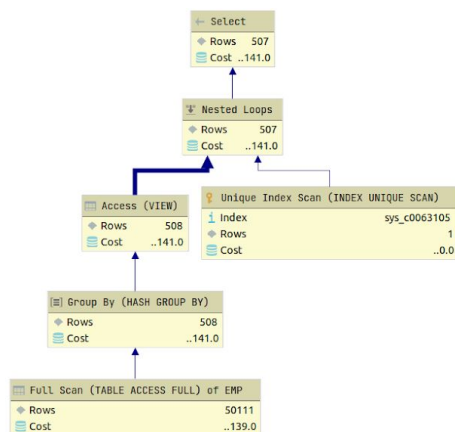
```
SELECT /*+ NO_USE_HASH(e d) */ d.deptno, AVG(e.sal)
```

```
FROM emp e, dept d
```

```
WHERE d.deptno = e.deptno
```

```
GROUP BY d.deptno;
```

Default=No_Use_Hash(e d)



Query2.sql - Query2.sql

Result 3

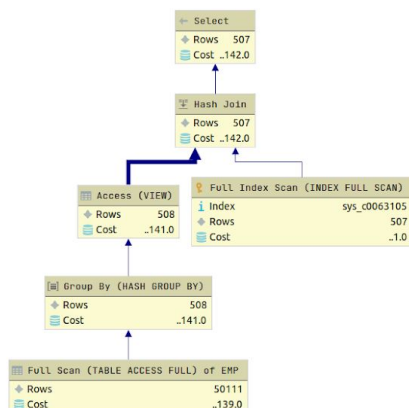
PLAN_TABLE_OUTPUT

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		507	21294	141 (3)	00:00:01
1	NESTED LOOPS		507	21294	141 (3)	00:00:01
2	VIEW	VW_GBC_5	508	19812	141 (3)	00:00:01
3	HASH GROUP BY		508	3048	141 (3)	00:00:01
4	TABLE ACCESS FULL	EMP	5011	293K	139 (1)	00:00:01
5	INDEX UNIQUE SCAN	SYS_C0063105	1	3	0 (0)	00:00:01

Predicate Information (identified by operation id):

5 - access("D"."DEPTNO"="ITEM_1")

Use_Hash(e d)



Result 9

PLAN_TABLE_OUTPUT

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		507	21294	142 (3)	00:00:01
1	HASH JOIN		507	21294	142 (3)	00:00:01
2	INDEX FULL SCAN	SYS_C0063105	507	1521	1 (0)	00:00:01
3	VIEW	VW_GBC_5	508	19812	141 (3)	00:00:01
4	HASH GROUP BY		508	3048	141 (3)	00:00:01
5	TABLE ACCESS FULL	EMP	5011	293K	139 (1)	00:00:01

Predicate Information (identified by operation id):

1 - access("D"."DEPTNO"="ITEM_1")

Query #3

Disabilitare il metodo hash join mediante l'uso di hint (/*+ NO_USE_HASH(e d) */)

```
SELECT /*+ NO_USE_HASH(e d) */ ename, job, sal, dname
```

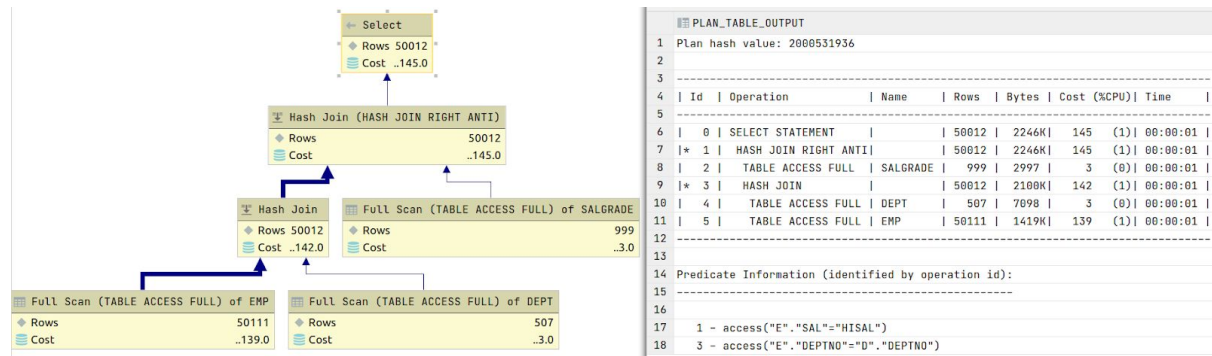
```
FROM emp e, dept d
```

```
WHERE e.deptno = d.deptno
```

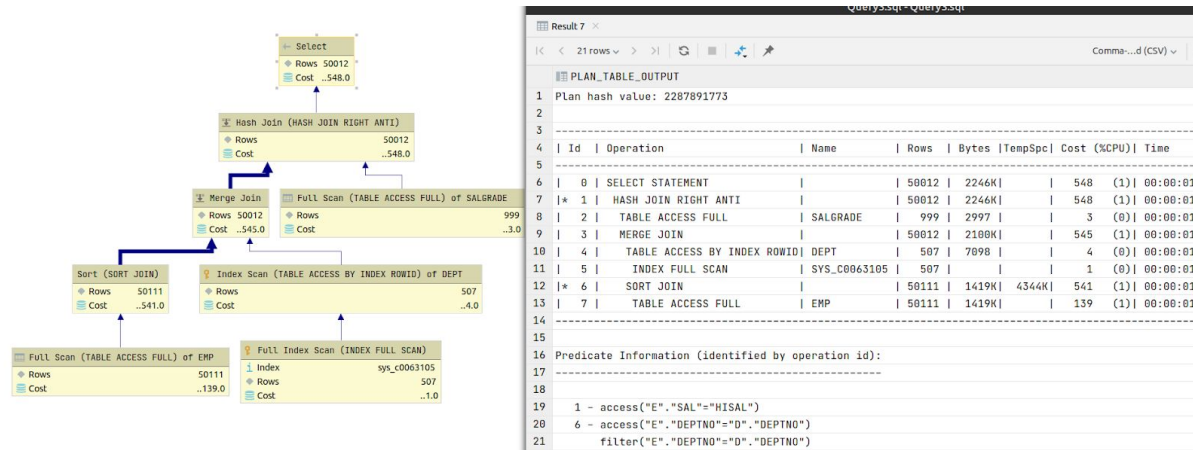
```
AND NOT EXISTS
```

```
(SELECT * FROM salgrade WHERE e.sal = hisal);
```

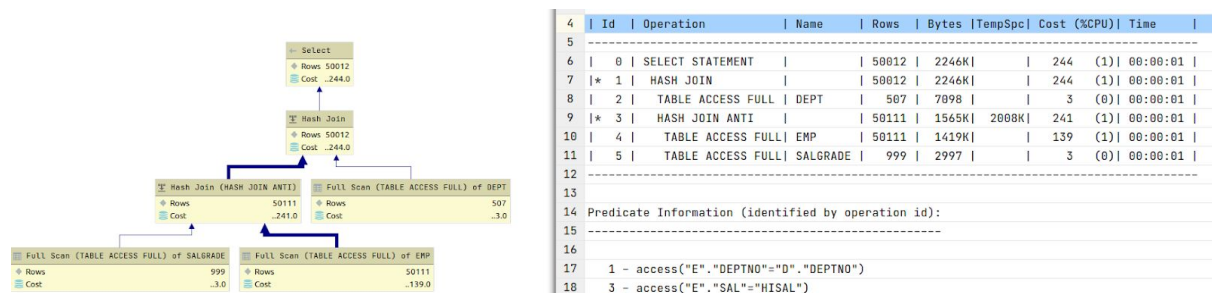
Default=Use_Hash(e d)



No_Use_Hash(e d)



/*+ ORDERED */(prima join tra emp e salgrade, e poi join con dept)



Query #4

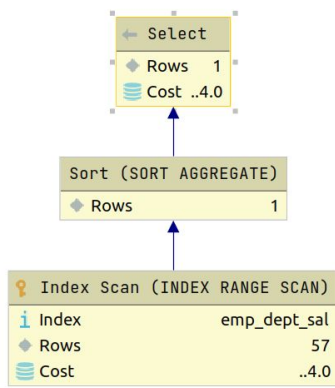
Si definiscano una o più strutture secondarie (indici) che permettano l'ottimizzazione della seguente query. Si analizzi con particolare attenzione il cambiamento nel piano di esecuzione creando due indici sugli attributi interessati dall'interrogazione.

```

select avg(e.sal)
from emp e
where e.deptno < 10 and
e.sal > 100 and e.sal < 200;

```

Default oppure deptno B-tree index oppure sal B-tree index



Result 17

15 rows

Comma...d (CSV)

PLAN_TABLE_OUTPUT

1

Plan hash value: 531818999

2

3

4

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	6	4 (0)	00:00:01
1	SORT AGGREGATE		1	6		
2	INDEX RANGE SCAN	EMP_DEPT_SAL	57	342	4 (0)	00:00:01

5

6

7

8

9

10

11

Predicate Information (identified by operation id):

12

13

14

2 - access("E"."SAL">100 AND "E"."DEPTNO"<10 AND "E"."SAL"<200)

15

filter("E"."SAL"<200 AND "E"."SAL">100)

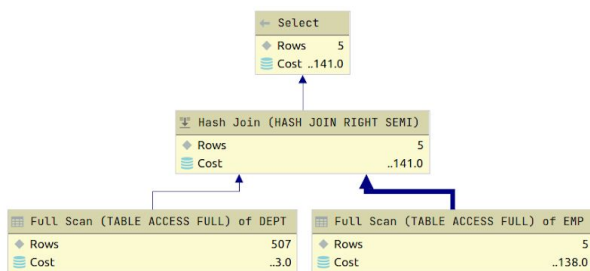
Query #5

Si definiscano una o più strutture secondarie (indici) che permettano l'ottimizzazione della seguente query:

```

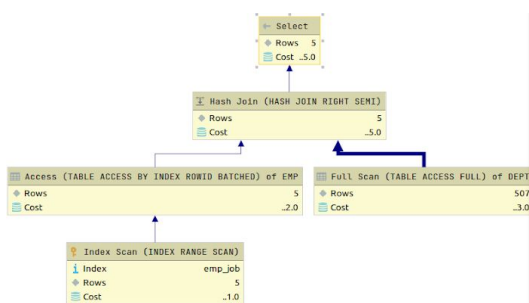
select dname
from dept
where deptno in (select deptno
from emp
where job = 'PHILOSOPHER');
  
```

Default o indice su dept(dname)



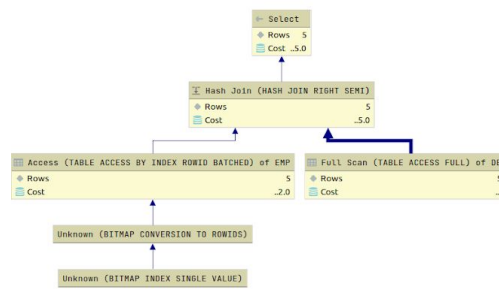
1	PLAN_TABLE_OUTPUT									
2										
3										
4	Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time			
5										
6	0	SELECT STATEMENT		5	130	141 (0)	00:00:01			
7	* 1	HASH JOIN RIGHT SEMI		5	130	141 (0)	00:00:01			
8	* 2	TABLE ACCESS FULL	EMP	5	60	138 (0)	00:00:01			
9	3	TABLE ACCESS FULL	DEPT	507	7098	3 (0)	00:00:01			
10										
11										
12	Predicate Information (identified by operation id):									
13										
14										
15	1 - access("DEPTNO"="DEPTNO")									
16	2 - filter("JOB"='PHILOSOPHER')									

Job B-tree index



1	Plan hash value: 262934622							
2								
3								
4	Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
5								
6	0	SELECT STATEMENT		5	130	5 (0)	00:00:01	
7	* 1	HASH JOIN RIGHT SEMI		5	130	5 (0)	00:00:01	
8	2	TABLE ACCESS BY INDEX ROWID BATCHED	EMP	5	60	2 (0)	00:00:01	
9	* 3	INDEX RANGE SCAN	EMP_JOB	5		1 (0)	00:00:01	
10	4	TABLE ACCESS FULL	DEPT	507	7098	3 (0)	00:00:01	
11								
12								
13	Predicate Information (identified by operation id):							
14								
15								
16	1 - access("DEPTNO"="DEPTNO")							
17	3 - access("JOB"='PHILOSOPHER')							

Job BITMAP index

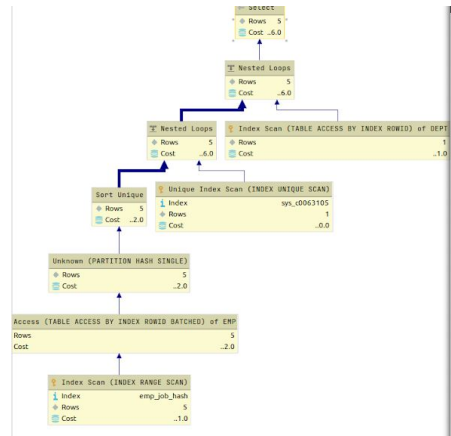


```

1 Plan hash value: 2288164076
2
3
4
5
6 Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |
7 --|-----|-----|-----|-----|-----|-----|-----|
8 0 | SELECT STATEMENT | | 5 | 130 | 5 (0)| 00:00:01 |
9 1 | HASH JOIN RIGHT SEMI | | 5 | 130 | 5 (0)| 00:00:01 |
10 2 | TABLE ACCESS BY INDEX ROWID BATCHED | EMP | 5 | 60 | 2 (0)| 00:00:01 |
11 3 | BITMAP CONVERSION TO ROWIDS | | | | | | |
12 4 | BITMAP INDEX SINGLE VALUE | EMP_JOB_BITMAP | | | | | |
13 5 | TABLE ACCESS FULL | DEPT | 507 | 7098 | 3 (0)| 00:00:01 |
14
15 Predicate Information (identified by operation id):
16
17 1 - access("DEPTNO"="DEPTNO")
18 4 - access("JOB"='PHILOSOPHER')

```

Job HASH index Nested Loop

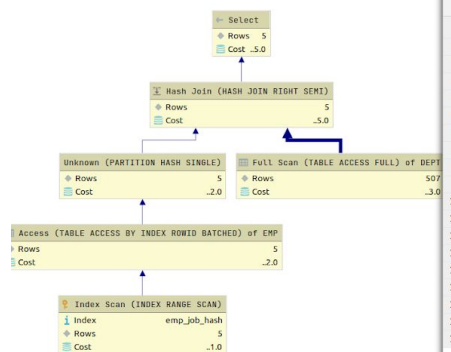


```

Result 44
21 rows
Comma...d (CSV)
PLAN_TABLE_OUTPUT
1 n hash value: 3293949767
2
3
4 Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time | Pstart| Pstop |
5 --|-----|-----|-----|-----|-----|-----|-----|
6 0 | SELECT STATEMENT | | 5 | 130 | 6 (17)| 00:00:01 | | |
7 1 | NESTED LOOPS | | 5 | 130 | 6 (17)| 00:00:01 | | |
8 2 | NESTED LOOPS | | 5 | 130 | 6 (17)| 00:00:01 | | |
9 3 | SORT UNIQUE | | 5 | 60 | 2 (0)| 00:00:01 | | |
10 4 | PARTITION HASH SINGLE | | 5 | 60 | 2 (0)| 00:00:01 | 1 | 1 |
11 5 | TABLE ACCESS BY INDEX ROWID BATCHED | EMP | 5 | 60 | 2 (0)| 00:00:01 | | |
12 6 | INDEX RANGE SCAN | EMP_JOB_HASH | 5 | | 1 (0)| 00:00:01 | 1 | 1 |
13 7 | INDEX UNIQUE SCAN | SYS_C0063105 | 1 | | 0 (0)| 00:00:01 | | |
14 8 | TABLE ACCESS BY INDEX ROWID | DEPT | 1 | 14 | 1 (0)| 00:00:01 | | |
15
16 Predicate Information (identified by operation id):
17
18 6 - access("JOB"='PHILOSOPHER')
19 7 - access("DEPTNO"="DEPTNO")

```

Job HASH index Hash Join

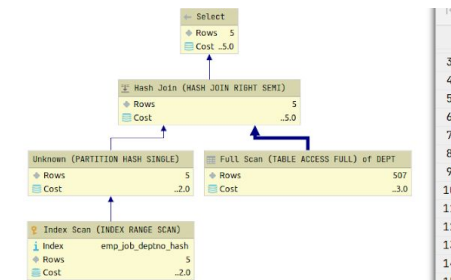


```

18 rows
Comma...d (CSV)
PLAN_TABLE_OUTPUT
1 Lan hash value: 4254969650
2
3
4 Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time | Pstart| Pstop |
5 --|-----|-----|-----|-----|-----|-----|-----|
6 0 | SELECT STATEMENT | | 5 | 130 | 5 (0)| 00:00:01 | | |
7 1 | HASH JOIN RIGHT SEMI | | 5 | 130 | 5 (0)| 00:00:01 | | |
8 2 | PARTITION HASH SINGLE | | 5 | 60 | 2 (0)| 00:00:01 | 1 | 1 |
9 3 | TABLE ACCESS BY INDEX ROWID BATCHED | EMP | 5 | 60 | 2 (0)| 00:00:01 | | |
10 4 | INDEX RANGE SCAN | EMP_JOB_HASH | 5 | | 1 (0)| 00:00:01 | 1 | 1 |
11 5 | TABLE ACCESS FULL | DEPT | 507 | 7098 | 3 (0)| 00:00:01 | | |
12
13 Predicate Information (identified by operation id):
14
15 1 - access("DEPTNO"="DEPTNO")
16 4 - access("JOB"='PHILOSOPHER')

```

Job,Deptno HASH index Hash Join

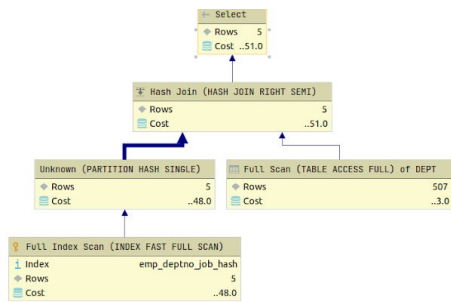


```

17 rows
Comma...d (CSV)
PLAN_TABLE_OUTPUT
3
4 Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time | Pstart| Pstop |
5 --|-----|-----|-----|-----|-----|-----|-----|
6 0 | SELECT STATEMENT | | 5 | 130 | 5 (0)| 00:00:01 | | |
7 1 | HASH JOIN RIGHT SEMI | | 5 | 130 | 5 (0)| 00:00:01 | | |
8 2 | PARTITION HASH SINGLE | | 5 | 60 | 2 (0)| 00:00:01 | 1 | 1 |
9 3 | INDEX RANGE SCAN | EMP_JOB_DEPTNO_HASH | 5 | 60 | 2 (0)| 00:00:01 | 1 | 1 |
10 4 | TABLE ACCESS FULL | DEPT | 507 | 7098 | 3 (0)| 00:00:01 | | |
11
12 Predicate Information (identified by operation id):
13
14 1 - access("DEPTNO"="DEPTNO")
15 3 - access("JOB"='PHILOSOPHER')

```

Deptno,Job HASH index Hash Join



PLAN_TABLE_OUTPUT

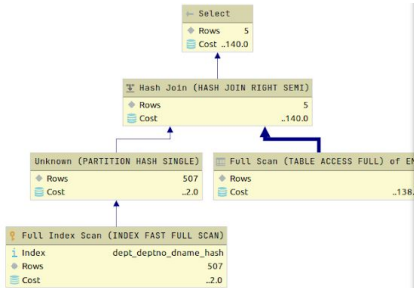
1 Plan hash value: 1140949081

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		5	130	51 (0)	00:00:01		
1	HASH JOIN RIGHT SEMI		5	130	51 (0)	00:00:01		
2	PARTITION HASH SINGLE		5	60	48 (0)	00:00:01	1	1
3	INDEX FAST FULL SCAN	EMP_DEPTNO_JOB_HASH	5	60	48 (0)	00:00:01	1	1
4	TABLE ACCESS FULL	DEPT	507	7098	3 (0)	00:00:01		

Predicate Information (identified by operation id):

1 - access("DEPTNO"="DEPTNO")
3 - filter("JOB"='PHILOSOPHER')

Deptno,Dname HASH index Hash Join



PLAN_TABLE_OUTPUT

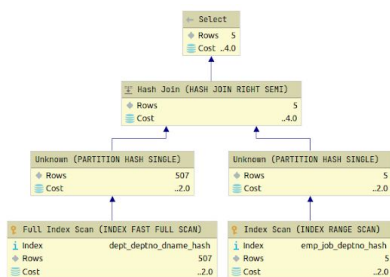
1 Plan hash value: 1763987074

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		5	130	140 (0)	00:00:01		
1	HASH JOIN RIGHT SEMI		5	130	140 (0)	00:00:01		
2	TABLE ACCESS FULL	EMP	5	60	138 (0)	00:00:01		
3	PARTITION HASH SINGLE		507	7098	2 (0)	00:00:01	1	1
4	INDEX FAST FULL SCAN	DEPT_DEPTNO_DNAME_HASH	507	7098	2 (0)	00:00:01	1	1

Predicate Information (identified by operation id):

1 - access("DEPTNO"="DEPTNO")
2 - filter("JOB"='PHILOSOPHER')

Job,Deptno HASH index + Deptno,Dname HASH index Hash Join



PLAN_TABLE_OUTPUT

1 Plan hash value: 1023227003

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		5	130	4 (0)	00:00:01		
1	HASH JOIN RIGHT SEMI		5	130	4 (0)	00:00:01		
2	PARTITION HASH SINGLE		5	60	2 (0)	00:00:01	1	1
3	INDEX RANGE SCAN	EMP_JOB_DEPTNO_HASH	5	60	2 (0)	00:00:01	1	1
4	PARTITION HASH SINGLE		507	7098	2 (0)	00:00:01	1	1
5	INDEX FAST FULL SCAN	DEPT_DEPTNO_DNAME_HASH	507	7098	2 (0)	00:00:01	1	1

Predicate Information (identified by operation id):

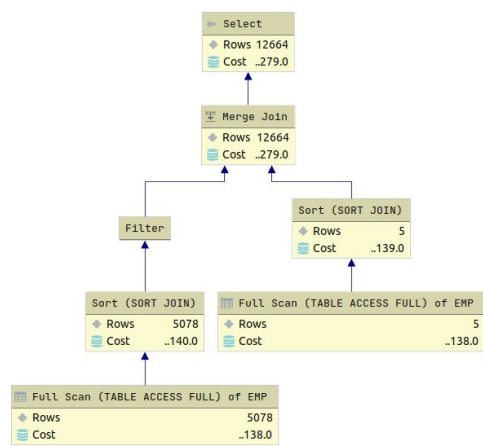
1 - access("DEPTNO"="DEPTNO")
3 - access("JOB"='PHILOSOPHER')

Query #6

Si definiscano una o più strutture secondarie (indici) che permettano l'ottimizzazione della seguente query (rimuovere eventuali indici già esistenti per confrontare le performance della query con e senza indici):

```
select e1.ename, e1.empno, e1.sal, e2.ename, e2.empno, e2.sal
from emp e1, emp e2
where e1.ename <> e2.ename and e1.sal < e2.sal
and e1.job = 'PHILOSOPHER' and e2.job = 'ENGINEER';
```

Default o SAL + Ename B-tree index



1 Plan hash value: 3733349388

2

3

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		12664	742K	279 (2)	00:00:01
1	MERGE JOIN		12664	742K	279 (2)	00:00:01
2	SORT JOIN		5	150	139 (1)	00:00:01
3	TABLE ACCESS FULL	EMP	5	150	138 (0)	00:00:01
4	FILTER					
5	SORT JOIN		5078	148K	140 (2)	00:00:01
6	TABLE ACCESS FULL	EMP	5078	148K	138 (0)	00:00:01

15 Predicate Information (identified by operation id):

16

17

18 3 - filter("E1"."JOB"='PHILOSOPHER')

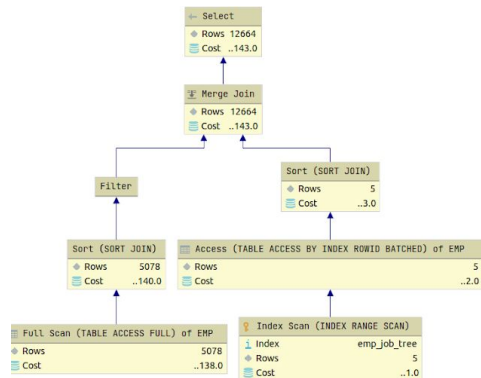
19 4 - filter("E1"."ENAME"<>"E2"."ENAME")

20 5 - access("E1"."SAL"<"E2"."SAL")

21 filter("E1"."SAL"<"E2"."SAL")

22 6 - filter("E2"."JOB"='ENGINEER')

Job B-Tree index



2

3

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		12664	742K	143 (3)	00:00:01
1	MERGE JOIN		12664	742K	143 (3)	00:00:01
2	SORT JOIN		5	150	3 (34)	00:00:01
3	TABLE ACCESS BY INDEX ROWID BATCHED	EMP	5	150	2 (0)	00:00:01
4	INDEX RANGE SCAN	EMP_JOB_TREE	5		1 (0)	00:00:01
5	FILTER					
6	SORT JOIN		5078	148K	140 (2)	00:00:01
7	TABLE ACCESS FULL	EMP	5078	148K	138 (0)	00:00:01

15 Predicate Information (identified by operation id):

16

17

18

19 4 - access("E1"."JOB"='PHILOSOPHER')

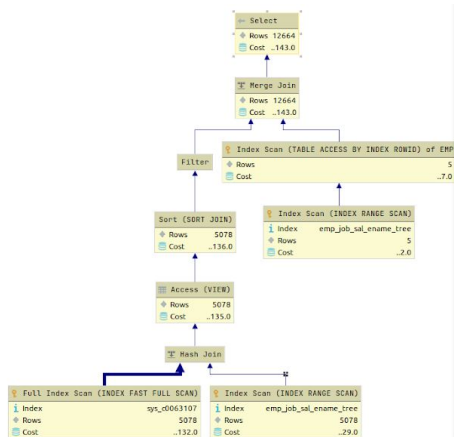
20 5 - filter("E1"."ENAME"<>"E2"."ENAME")

21 6 - access("E1"."SAL"<"E2"."SAL")

22 filter("E1"."SAL"<"E2"."SAL")

23 7 - filter("E2"."JOB"='ENGINEER')

Job,Sal,Ename B-Tree index



Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		12664	742K	143 (2)	00:00:01
1	MERGE JOIN		12664	742K	143 (2)	00:00:01
2	TABLE ACCESS BY INDEX ROWID	EMP	5	150	7 (0)	00:00:01
3	INDEX RANGE SCAN	EMP_JOB_SAL_ENAME_TREE	5		2 (0)	00:00:01
4	FILTER					
5	SORT JOIN		5078	148K	136 (2)	00:00:01
6	VIEW	index_join\$_002	5078	148K	135 (1)	00:00:01
7	HASH JOIN					
8	INDEX RANGE SCAN	EMP_JOB_SAL_ENAME_TREE	5078	148K	29 (0)	00:00:01
9	INDEX FAST FULL SCAN	SYS_C0063107	5078	148K	132 (1)	00:00:01

15 Predicate Information (identified by operation id):

16

17

18

19

20

21 3 - access("E1"."JOB"='PHILOSOPHER')

22 4 - filter("E1"."ENAME"<>"E2"."ENAME")

23 5 - access("E1"."SAL"<"E2"."SAL")

24 filter("E1"."SAL"<"E2"."SAL")

25 6 - filter("E2"."JOB"='ENGINEER')

26 7 - access(ROWID=ROWID)

27 8 - access("E2"."JOB"='ENGINEER')